

LINUX KERNEL

HACKING:

A CRASH COURSE

HELLO, MY NAME IS GEORGI

@GeorgiCodes

AGENDA

History of Linux

Kernel architecture

How to build a kernel module

How to build a device driver



debian



ubuntu.



CentOS



Core OS



kernel v4.0

8000 developers

800 companies

15 million lines of code

10 patches 7/365

2-3 months new release



I'm doing a (free)
operating system (just a
hobby, won't be big and
professional like gnu) ...

- Linus Torvalds 1991

TIMELINE

1991: Linus Torvalds creates kernel prototype

1994: Linux version 1.0.0 released

Mid 1990s: Lots of Linux distributions

1996: Tux was born

1994-1997: Linux gets mainstream press

1998: Support from Google, Oracle, Intel & Netscape

.....

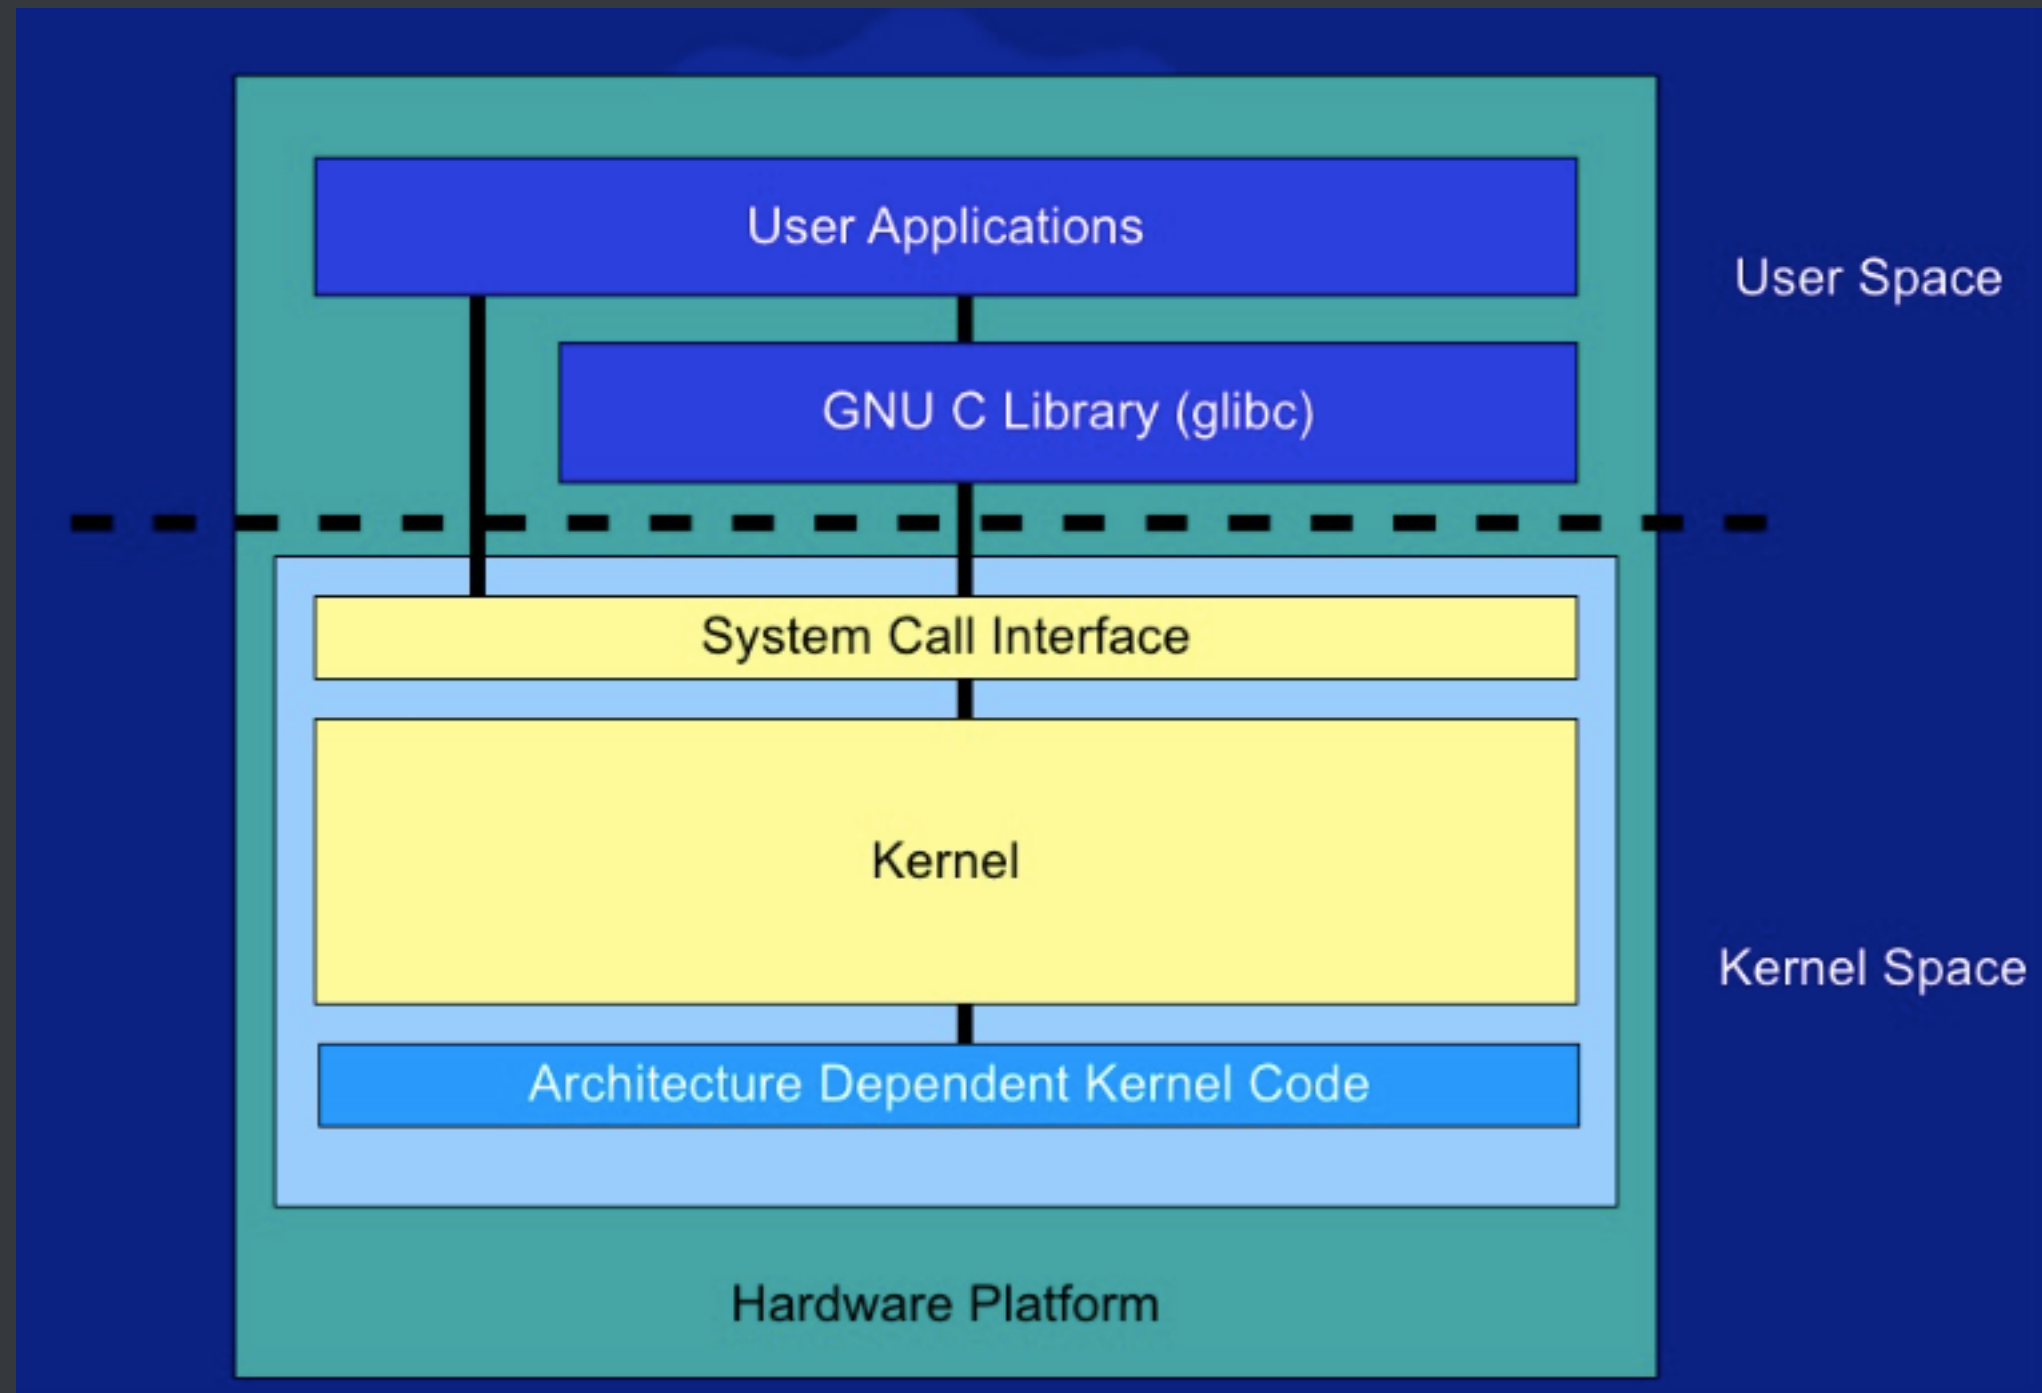
2015: Linux kernel version 4.0 released



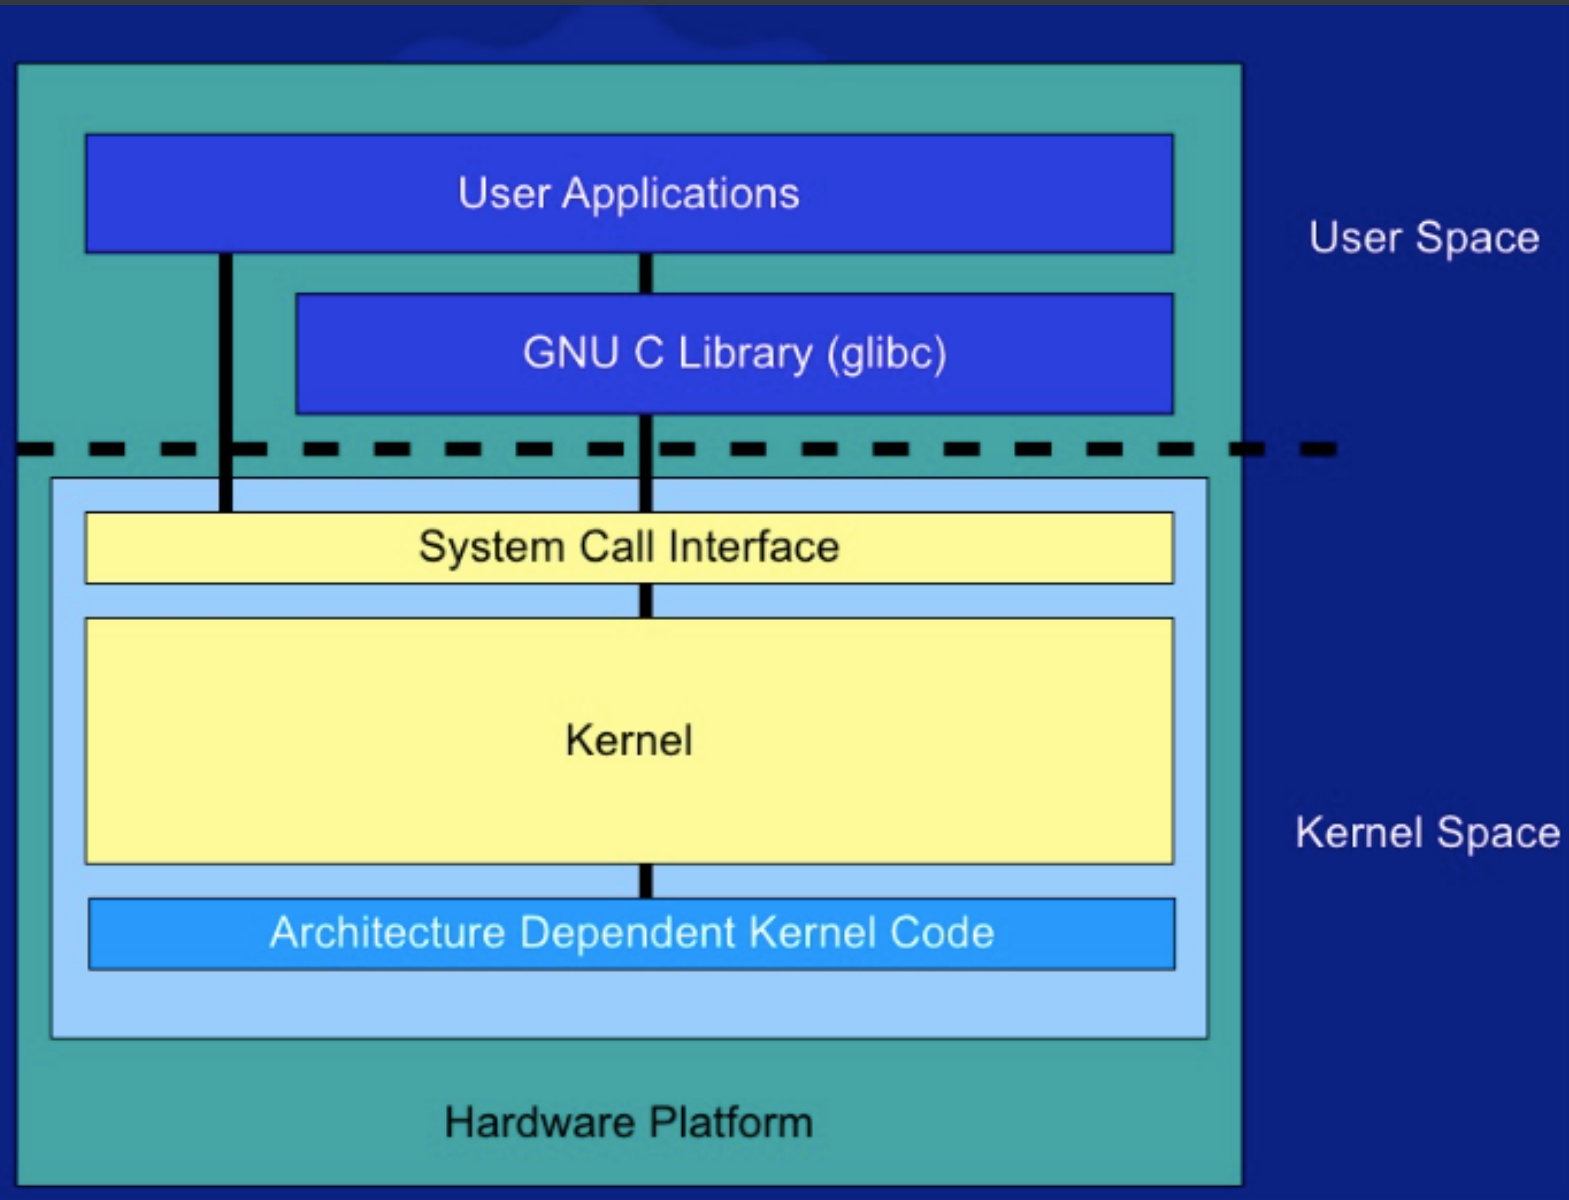
“... the **kernel** is a computer program that manages I/O (input/output) requests from software, and translates them into data processing instructions for the central processing unit and other electronic components of a computer.”

– **Wikipedia**

KERNEL ARCHITECTURE



KERNEL VS. USER SPACE



* User space restricts user programs so that they can't accidentally mess with the system.

* Kernel space is privileged and has full access to memory and resources.

WHO USES THE KERNEL?

The kernel provides a way for other programs to use the hardware via system calls.

- * The kernel is not designed for direct human consumption (no UI).
- * The kernel's users are other programs.

SYSTEM CALLS

- * there are a few hundred sys calls with functions like `read()`, `write()`, `open()`
- * When a system call is executed, the arguments are passed from `user` space to `kernel` space.
- * A `user` process becomes a `kernel` process when it executes a system call.

```
#include <fcntl.h>
int main()
{
    int fd, count; char buf[1000];
    fd=open("mydata", O_RDONLY);
    count = read(fd, buf, 1000);
    write(1, buf, count);
    close(fd);
}
```

WHERE DOES KERNEL CODE EXECUTE?

The kernel has **two** entry points:

1. When a running process/application that makes a **system call**.

2. Responding to a **hardware interrupt**

- A key was pressed
- A network packet just arrived
- A time just ticked

/PROC

- * window of communication between the kernel and user space
- * dynamically generated files provide info on running system

WHAT ARE KERNEL MODULES?

- * programs written in C built against the Linux kernel source tree
- * run in kernel space
- * core of kernel remains small where modules can be loaded and unloaded as required

To Build:

- * need kernel source tree, gcc and make
- * run the same version of kernel you built module with

```
#include <linux/module.h>    // included for all kernel modules
#include <linux/kernel.h>    // included for KERN_DEUBG
#include <linux/init.h>      // included for __init and __exit macros

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Georgi");
MODULE_DESCRIPTION("A Simple Hello World module");

static int __init hello(void)
{
    printk(KERN_DEBUG ">>> Hello world! <<<\n");
    return 0;    // Non-zero return means that the module couldn't be loaded.
}

static void __exit goodbye(void)
{
    printk(KERN_DEBUG ">>> Goodbye world! <<<\n");
}

module_init(hello);
module_exit(goodbye);
```

BUILD YOUR KERNEL MODULE

```
ifneq ($(KERNELRELEASE),)
obj-m    := hello.o

else
    KDIR ?= /lib/modules/`uname -r`/build

default:
    $(MAKE) -C $(KDIR) M=$$PWD
endif
```

COMMANDS

Loading + Unloading

- * `insmod hello.ko` - loads module
- * `rmmmod hello.ko` - unloads module

View kernel logs

- * `dmesg`

Use modprobe to manage dependencies

- * `Copy hello.ko into /lib/modules/$KERNEL_VERSION`
- * `depmod`
- * `modprobe hello` - loads module
- * `modprobe -r hello` - unloads module
- * `modinfo` - tells you info about the module

DEMO!

HOW IS KERNEL PROGRAMMING DIFFERENT?

1. kernel has no standard C headers and libraries
2. no memory protection!
3. a single big namespace
4. always multi-threaded

TYPES OF DEVICE DRIVERS

Char

- * reads/writes character by character to the device
- * operates in a blocking mode
- * stream of bytes

Block

- * reads/writes large amounts of data block by block.
- * operates in a non-blocking mode

Network device

- * Exchange data over network
- * Understands packets and connections

USB device

LET'S BUILD A DEVICE DRIVER!

What will our `char` device driver do?

It will respond with *"Hello Code PaLOUsa!"* when its read from.

Steps

1. Write the code
2. Build and load our module
3. Create a device file

.....

```
int init_module(void)
{
    major = register_chrdev(0, DEVICE_NAME, &fops);

    printk(KERN_INFO ">>> I was assigned major number %d. <<<\n", major);
    printk(KERN_INFO ">>> Run 'mknod /dev/%s c %d 0'. <<<\n", DEVICE_NAME, major);

    return 0;
}

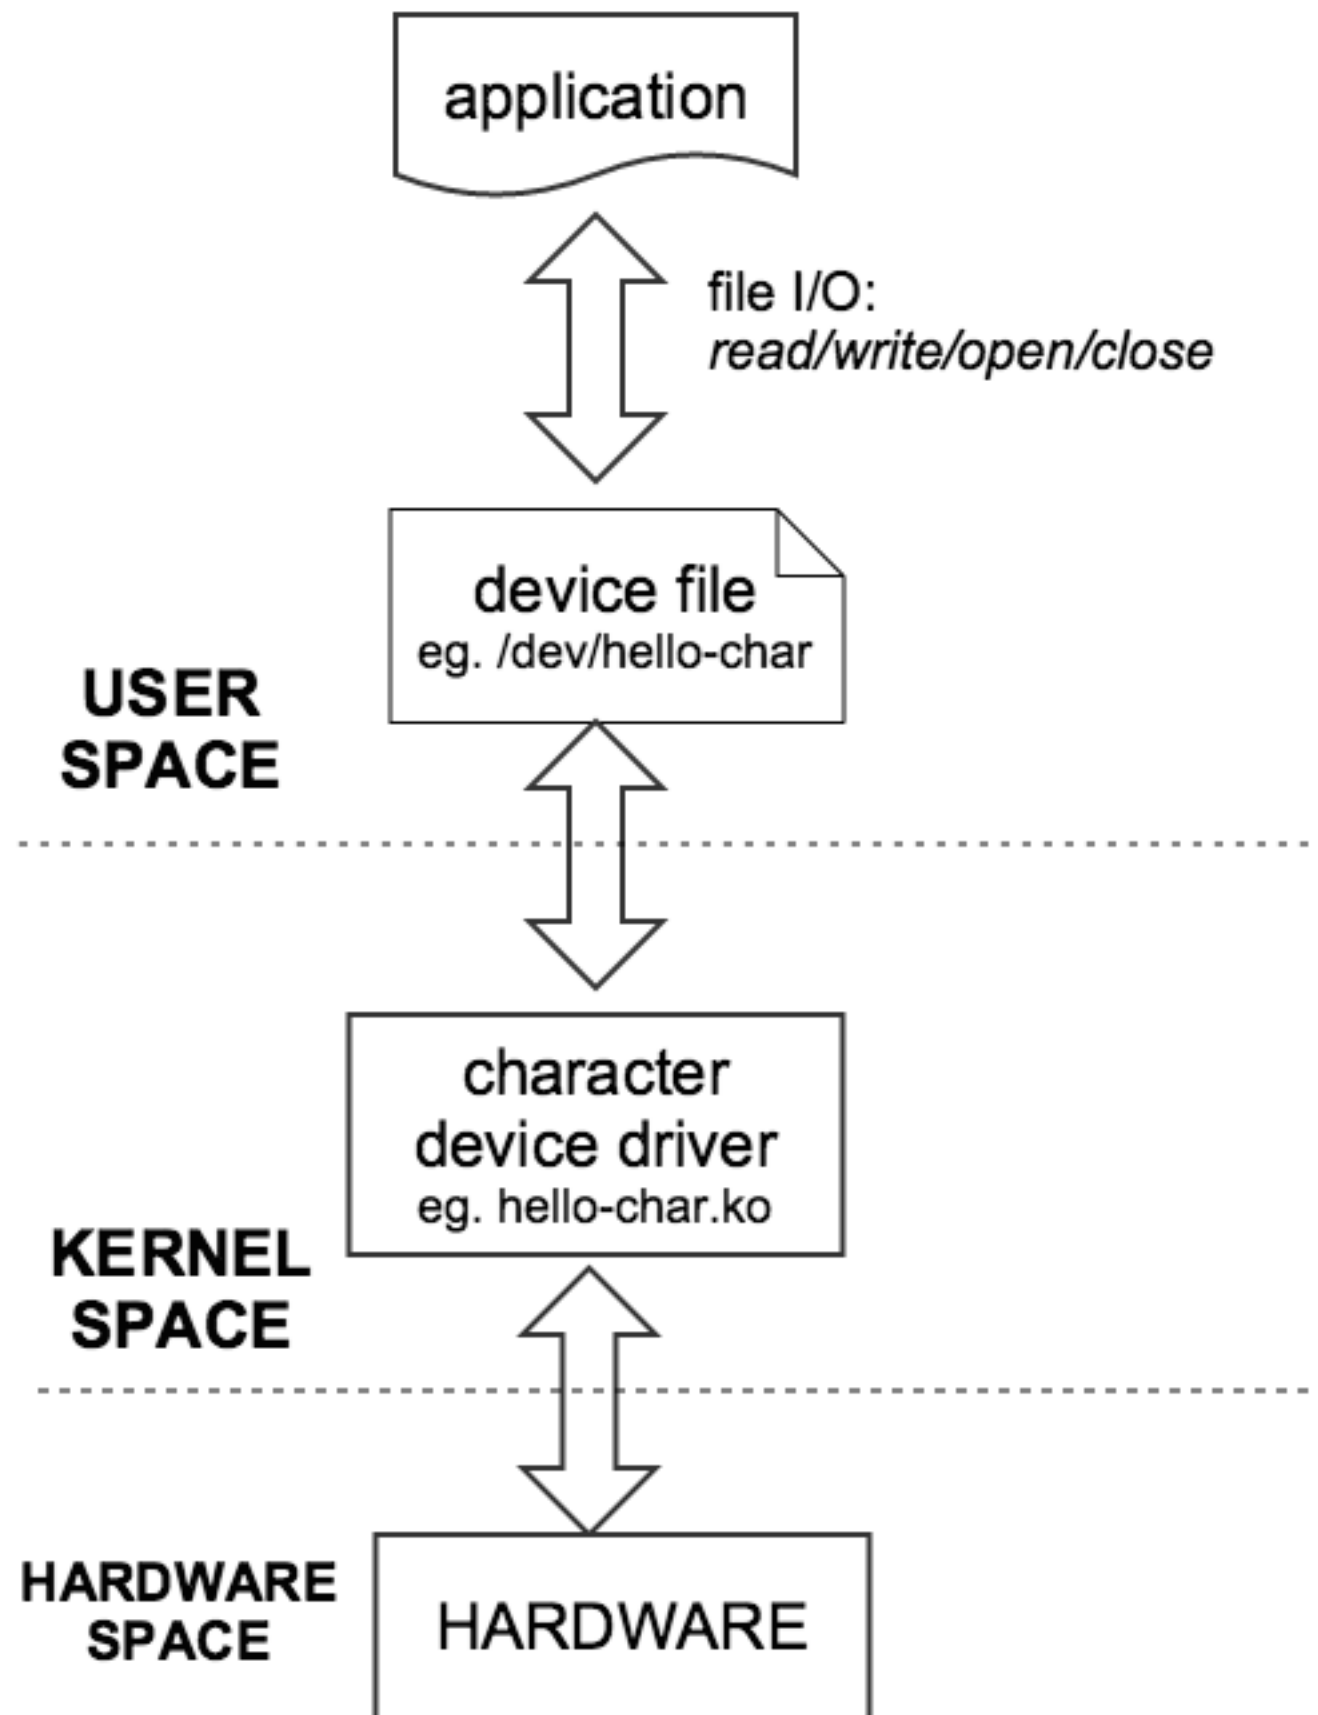
void cleanup_module(void)
{
    unregister_chrdev(major, DEVICE_NAME);
}
```

.....

```
static struct file_operations fops = {  
    .read = read_dev,  
    .write = write_dev,  
    .open = open_dev,  
    .release = close_dev  
};
```

* interaction is through system calls: `open()`,
`close()`, `read()`, `write()`
* “ops” struct pattern that allows you to fill in
behaviors via `callbacks`

DEVICE FILES



- * a device file is how a user program can access the physical device which lives all the way in kernel space.

- * device files live in the `/dev` directory

- * create device file:
`mknod /dev/hello-char c 250 0`

OUR DEVICE DRIVER IN ACTION!

Reading

```
cat /dev/hello-char
```

this will call the `read_dev()` function

DEMO!

LEVELLING UP YOUR LINUX SKILLS

- * install native Linux and use it!
- * configure and build your own kernel
- * write your own kernel module and/or device driver
- * do the Eudryptula Challenge

code + links + slides:

georgi.io/kernel-talk

[BITLY.IS/HIRING](https://bitly.is/hiring)